

# Comprehensive SBERT Fine-tuning Process for Bengali Text Classification

## Overview

Based on the latest SBERT v5+ documentation, this process covers fine-tuning a SentenceTransformer model for your 10k Bengali customer support classification dataset using the modern

`SentenceTransformerTrainer` approach.

## Step 1: Model Selection and Initialization

### Option A: Start with Bengali-Specific Model (Recommended)

```
python

from sentence_transformers import SentenceTransformer

# Best for Bengali: Start with BanglaBERT base
model = SentenceTransformer("csebuetnlp/banglabert")
```

### Option B: Start with Multilingual Model

```
python

# Alternative: Multilingual model with Bengali support
model = SentenceTransformer("paraphrase-multilingual-mpnet-base-v2")
```

### Option C: Build from Transformer Base

```
python

from sentence_transformers import SentenceTransformer, models

# Create from scratch using Bengali transformer
word_embedding_model = models.Transformer("csebuetnlp/banglabert", max_seq_length=256)
pooling_model = models.Pooling(word_embedding_model.get_word_embedding_dimension())
model = SentenceTransformer(modules=[word_embedding_model, pooling_model])
```

## Step 2: Dataset Preparation

### Format 1: Using SoftmaxLoss (Classification-Specific)

```
python

from datasets import Dataset
import pandas as pd

# Convert your data to the required format
data = {
    "sentence1": [], # First sentence from pair (can be empty string)
    "sentence2": [], # Your actual Bengali text
    "label": [] # Integer labels (0, 1, 2, ..., n_classes-1)
}

# Example with your Bengali data
train_data = []
for question, tag in your_data:
    # Create label mapping
    label_map = {
        'goodbye': 0,
        'greetings': 1,
        'agent_calling': 2,
        'repeat_again': 3,
        'namjari_process': 4,
        # ... etc for all your tags
    }

    train_data.append({
        'sentence1': "", # Empty string for single sentence classification
        'sentence2': question, # Your Bengali text
        'label': label_map[tag]
    })

train_dataset = Dataset.from_list(train_data)
```

### Format 2: Using MultipleNegativesRankingLoss (Contrastive Approach)

```
python
```

```
# Convert to (anchor, positive) pairs by grouping same classes
```

```
from collections import defaultdict
```

```
class_groups = defaultdict(list)
```

```
for question, tag in your_data:
```

```
    class_groups[tag].append(question)
```

```
pairs_data = []
```

```
for tag, questions in class_groups.items():
```

```
    # Create positive pairs within same class
```

```
    for i, q1 in enumerate(questions):
```

```
        for j, q2 in enumerate(questions[i+1:], i+1):
```

```
            pairs_data.append({
```

```
                'anchor': q1,
```

```
                'positive': q2
```

```
            })
```

```
train_dataset = Dataset.from_list(pairs_data)
```

## Step 3: Loss Function Selection

### Option A: SoftmaxLoss (Direct Classification)

```
python
```

```
from sentence_transformers.losses import SoftmaxLoss
```

```
num_labels = len(set(your_label_values)) # Number of unique classes
```

```
embedding_dim = model.get_sentence_embedding_dimension()
```

```
loss = SoftmaxLoss(
```

```
    model=model,
```

```
    sentence_embedding_dimension=embedding_dim,
```

```
    num_labels=num_labels,
```

```
    concatenation_sent_rep=True,
```

```
    concatenation_sent_difference=True,
```

```
    concatenation_sent_multiplication=False
```

```
)
```

### Option B: MultipleNegativesRankingLoss (Contrastive Learning)

```
python
```

```
from sentence_transformers.losses import MultipleNegativesRankingLoss
```

```
loss = MultipleNegativesRankingLoss(model=model, scale=20.0)
```

## Option C: CoSENTLoss (Modern Approach with Similarity Scores)

If you can generate similarity scores between pairs:

```
python
```

```
from sentence_transformers.losses import CoSENTLoss
```

```
loss = CoSENTLoss(model=model, scale=20.0)
```

## Step 4: Training Arguments Configuration

```
python
```

```
from sentence_transformers import SentenceTransformerTrainingArguments
from sentence_transformers.training_args import BatchSamplers

args = SentenceTransformerTrainingArguments(
    # Required
    output_dir="/results/bengali-sbert-finetuned",

    # Core training parameters
    num_train_epochs=3, # Start with 3, adjust based on validation
    per_device_train_batch_size=16, # Adjust based on GPU memory
    per_device_eval_batch_size=16,
    learning_rate=2e-5, # Standard for transformer fine-tuning
    warmup_ratio=0.1,

    # Performance optimization
    fp16=True, # Use if your GPU supports it
    bf16=False, # Use if you have newer GPU (A100, etc.)
    gradient_checkpointing=True, # Save memory

    # For SoftmaxLoss, ensure each batch has examples from each class
    batch_sampler=BatchSamplers.GROUP_BY_LABEL if using_softmax else BatchSamplers.NO_DUPLICATES,

    # Evaluation and saving
    eval_strategy="steps",
    eval_steps=100,
    save_strategy="steps",
    save_steps=100,
    save_total_limit=2,
    load_best_model_at_end=True,
    metric_for_best_model="eval_loss",

    # Logging
    logging_steps=50,
    report_to=["tensorboard"], # or ["wandb"] if you use Weights & Biases
    run_name="bengali-sbert-classification",

    # Optional: For prompts (if using instruction-based training)
    # prompts="Classify this Bengali customer service query: ",
)
```

## Step 5: Evaluation Setup

```
python

from sentence_transformers.evaluation import EmbeddingSimilarityEvaluator, LabelAccuracyEvaluator

# For classification tasks
evaluator = LabelAccuracyEvaluator(
    sentences=eval_sentences,
    labels=eval_labels,
    name="bengali-classification-eval"
)
```

## Step 6: Training Execution

```
python

from sentence_transformers import SentenceTransformerTrainer

trainer = SentenceTransformerTrainer(
    model=model,
    args=args,
    train_dataset=train_dataset,
    eval_dataset=eval_dataset, # Optional
    loss=loss,
    evaluator=evaluator, # Optional
)

# Start training
trainer.train()

# Save the final model
model.save_pretrained("./bengali-sbert-final")
```

## Step 7: Advanced Configurations

### A. Data Augmentation (Recommended for 10k samples)

```
python
```

*# Create synthetic examples through back-translation or paraphrasing*

```
from transformers import pipeline
```

*# Example: Use mT5 for Bengali paraphrasing*

```
paraphraser = pipeline("text2text-generation", model="google/mt5-base")
```

```
def augment_data(text):
```

```
    prompt = f"paraphrase: {text}"
```

```
    return paraphraser(prompt, max_length=100)[0]['generated_text']
```

*# Apply to your dataset*

```
augmented_data = []
```

```
for item in original_data:
```

```
    augmented_data.append(item) # Original
```

```
    augmented_text = augment_data(item['sentence2'])
```

```
    augmented_data.append(**item, 'sentence2': augmented_text)
```

## B. Hyperparameter Optimization

```
python
```

```

from sentence_transformers.trainer import SentenceTransformerTrainer

def model_init():
    return SentenceTransformer("csebuetnlp/banglabert")

def hp_space(trial):
    return {
        "learning_rate": trial.suggest_float("learning_rate", 1e-5, 5e-5, log=True),
        "per_device_train_batch_size": trial.suggest_categorical("per_device_train_batch_size", [8, 16, 32]),
        "num_train_epochs": trial.suggest_int("num_train_epochs", 2, 5),
    }

trainer = SentenceTransformerTrainer(
    model_init=model_init,
    args=args,
    train_dataset=train_dataset,
    eval_dataset=eval_dataset,
    loss=loss,
)

# Run hyperparameter optimization
best_trial = trainer.hyperparameter_search(
    hp_space=hp_space,
    compute_objective=lambda metrics: metrics["eval_loss"],
    n_trials=20,
    direction="minimize",
)

```

## C. Multi-dataset Training (If you have additional datasets)

```
python
```



```
# Combine multiple datasets
train_datasets = {
    "main": main_train_dataset,
    "auxiliary": auxiliary_dataset, # Could be translated NLI data
}

# Different loss functions for different datasets
losses = {
    "main": SoftmaxLoss(model, embedding_dim, num_labels),
    "auxiliary": MultipleNegativesRankingLoss(model),
}

trainer = SentenceTransformerTrainer(
    model=model,
    train_dataset=train_datasets,
    loss=losses,
    args=args,
)
```

## Step 8: Model Usage After Training

```
python
```

```
# Load your fine-tuned model
model = SentenceTransformer("./bengali-sbert-final")

# For classification, you'll need to add a classifier on top
from sklearn.linear_model import LogisticRegression
import numpy as np

# Generate embeddings for your training data
train_embeddings = model.encode(train_texts)
train_labels = your_train_labels

# Train classifier
classifier = LogisticRegression()
classifier.fit(train_embeddings, train_labels)

# For prediction
def classify_text(text):
    embedding = model.encode([text])
    prediction = classifier.predict(embedding)[0]
    confidence = classifier.predict_proba(embedding)[0].max()
    return prediction, confidence

# Example usage
text = "আমি নামজারি করতে চাই"
label, confidence = classify_text(text)
print(f"Predicted: {label}, Confidence: {confidence:.3f}")
```

## Key Considerations for Bengali

1. **Tokenization:** BanglaBERT handles Bengali tokenization properly
2. **Sequence Length:** Monitor your text lengths, adjust `max_seq_length` if needed
3. **Mixed Script:** Handle mixed Bengali-English text if present in your data
4. **Class Imbalance:** Use appropriate sampling strategies if classes are imbalanced

## Performance Tips

1. **Batch Size:** Start with 16, increase if GPU memory allows
2. **Learning Rate:** 2e-5 is generally good, but try 1e-5 for stability
3. **Epochs:** 3-5 epochs usually sufficient for fine-tuning
4. **Validation:** Always use a held-out validation set

5. **Early Stopping:** Monitor validation loss to prevent overfitting

## Expected Results

With 10k examples and proper fine-tuning:

- Classification accuracy: 85-92% (depending on task complexity)
- Training time: 2-4 hours on single GPU
- Model size: ~440MB for BanglaBERT base